

Four audiences, one month


The same August data, distributed four ways. What separates these documents is not formatting – it is what each one is forbidden to contain.

month 2026-08 model claude-opus-5 runs 2 PRs merged 7 engineers 1


LAYER 1 → LAYER 2 · DETERMINISTIC

Seven merged pull requests become four

Every PR is classified by the files it touches, before anything counts it. No model runs here, so the split is not arguable – the same file list always produces the same answer.



COUNTED



CLASSIFIED OUT



Both numbers appear on every page that shows either. Nothing is hidden – the split sums back to the raw count.

LAYER 2 · TEAM

Four of seven metrics have no answer

A missing source produces a blank, never a zero. A change-failure rate of 0 because nobody wired up an incident tracker reads as good news; that is the failure mode this avoids.

METRIC	VALUE	BASIS
Deployment frequency	<u>insufficient evidence</u>	Needs a deploy feed. A merge is not a deploy.
Lead time for changes	0 h	Median PR open → merge, n=7. No deploy feed, so this is merge time, not deploy time.
Change failure rate	<u>insufficient evidence</u>	Needs an incident feed. Revert rate is a weak proxy.

METRIC	VALUE	BASIS
Time to restore	<u>insufficient evidence</u>	Needs incidents carrying <code>resolved_at</code> .
Review latency	<u>insufficient evidence</u>	Needs a review by someone other than the author — none of the 7 PRs inspected had one.
Work in progress	0.5 PRs	Median concurrently-open PRs. Proxy: no board feed configured.
Carryover	0 %	0 of 7 merged PRs were opened in an earlier month.

The review-latency line is the one worth reading twice. It is not a tooling gap — it is the finding: seven pull requests merged, none reviewed by anyone but their author.

LAYER 4 · POLICY

Who receives what

Four documents from one dataset. Each card leads with what that audience gets, and ends with what it is structurally denied.

RELEASED FIRST

The engineer

THEIR OWN PAGE

- Their own throughput, type split, reviews, rework
- The full interpretation, both readings
- Claims the gate deleted, and why
- How to file a correction

NEVER CONTAINS

Anyone else's data. No comparison exists, because none was computed.

The manager

1:1 AGENDA

- Everything, framed as questions to ask
- Contested claims marked as contested
- What the data cannot show, listed first

NEVER CONTAINS

Any ordering that means anything. It is an agenda, not a verdict.

Squad leads

OWN SQUAD + AGGREGATE

- Their squad's throughput by type
- The team aggregate for context

NEVER CONTAINS

Another squad's individuals. Rows are alphabetical.

The founder

TEAM METRICS AND RISK

- DORA and flow at team level
- Risk, including what cannot be measured

NEVER CONTAINS

Any individual data at all.

A fifth document is never generated. There is no ranked list — not suppressed at render time, not access-controlled. Layer 2 computes no score, Layer 3's schema has no field for one, and the gate refuses any list of engineers carrying an ordinal. Three independent checks, each of which fails the build rather than shipping a ranking.

“Engineer first” is enforced, not intended. The manager-facing pages *cannot be written* until 24 hours after the engineer pages are stamped. Running them early exits 3.

LAYER 3 → GATE

What survived the gate

The model ran twice. Every claim below cites a real diff, because claims that cited nothing checkable were deleted in code before any human saw them.

0

CLAIMS DROPPED
FOR NO CITATION

0

SCORE KEYS
STRIPPED

0

COMPARATIVE
PHRASES REMOVED

32

CLAIM-LEVEL
DIFFERENCES

Both runs independently raised the same work.

PR 1

PR 3

PR 4

PR 6

PR 7

PR 2 — raised by one run only

Where the two runs differ is in wording and emphasis, not in which work was worth describing. Claim-level comparison is textual, so a restatement can appear as a difference — which is why the PR-level signal above leads, and nothing is ever averaged.

PR 1 chose the paginated `pulls` list endpoint over the Search API specifically to avoid a silent undercount at high PR volume, and documented the reasoning in the module docstring.

PR #1 • diff

```
Uses the paginated `/repos/{owner}/{repo}/pulls` list endpoint (sorted by
created/updated) rather than the Se...
```

PR 1 handles GitHub rate limiting with a bounded retry that honours `Retry-After` and falls back to exponential backoff.

PR #1 • diff

```
if exc.code in (403, 429) and attempt < 5:
    wait = int(exc.headers.get("Retry-After", "0")) or...
```

PR 6 adds operator-facing troubleshooting material mapping failure symptoms to causes and fixes, which is written for someone other than the author.

PR #6 • diff

```
| Run is red on **Commit and push results** | Another run pushed first. The step
fetches, rebases and retries...
```

The same month, read two ways

Layer 3 cannot return one without the other. Both are readings of identical evidence; neither is the verdict.

MOST FAVOURABLE	LEAST FAVOURABLE
Engineer A shipped a working system end to end in one month and then, having seen what it produced, replaced its central premise. The engineering along the way is careful about failure modes others would leave implicit: the Search API cap, the concurrent-push race, the difference between “no data” and “zero” in the trend table.	Seven PRs merged with zero external review and a median time to merge of 0.0 hours, including a 22,317-line rewrite that discards the architecture of PR 1 four PRs later. Substantive throughput is 4 of 7 merged PRs, and the most load-bearing claims are visible only as prose in the README.

What this data cannot show

Sixteen gaps were returned. These are the ones a manager would most need before opening their mouth in a 1:1.

- ? Whether the redesign was self-directed, requested by someone, or a response to feedback received outside a pull request.
- ? Whether the absence of reviews reflects a solo repo with no available reviewers, a choice not to seek review, or a team norm.
- ? Whether the 0.0 h median time to merge reflects a self-merge policy, a solo repo, or something else.
- ? Any design work, scoping conversation, mentoring, or incident response that happened outside a pull request.
- ? Whether this repository has any users or collaborators beyond the author — contributors: 1 measures PR authorship, not usage.

The interpretation also flagged that “123 tests” appeared only as README prose with no test files in the visible diff, and questioned whether the pipeline had ever run with a real API key. Both were fair at the time it wrote them. The second is no longer true — this report is the evidence.

Generated by the four-layer pipeline at main. Layers 1, 2 and the gate are code; Layer 3 is the model; Layer 4 is policy.

Source documents: reports/balajilandge-Engineering-Metrics/2026-08/ — engineers/balajilandge.md, em.md, squads/team.md, founder.md. Gate audit: data/.../2026-08.interpretation.json.